



ECSE420 Parallel Computing
Lab 1 Logic Gates Simulation
Group 43
Hanwen Wang 260778557
Mai Zeng 260782174

1. Write code that simulates logic gate sequentially:

Code:

The code for sequential part is running on CPU only and because the program is only running on CPU then we just running the program on macOS system terminal, allocating the memory on RAM using malloc.

In the main, we get the input file pointer and output file pointer and then pass these parameters along with the length of input file length (the total number of lines in the input file) to the parseFile function. The reason we multiply the input file length by 6 is because there are 6 characters for each line in the input file. The further details we will talk later.

```
85 int main(int argc, char* argv[])
86 {
87     if ( argc < 4)
88     {
89         printf("You must enter 3 input files!\n");
90         exit(1);
91     }
92
93     // argv[1] : input_file_path
94     // argv[2] : input_file_length
95     // argv[3] : output_file_path
96
97     char* fileName = argv[1];
98     FILE* input = fopen(fileName, "r");
99     if (input == NULL)
100     {
101         exit(EXIT_FAILURE);
102     }
103     char* outputFileName = argv[3];
104     FILE* output = fopen(outputFileName, "w");
105     if (output == NULL)
106     {
107         exit(EXIT_FAILURE);
108     }
109     parseFile(input, atoi(argv[2])*6, output);
110
111     fclose(input);
112     fclose(output);
113
114     return 0;
115 }
```

Figure 1.1: CPU Code Snippet Main

The code snippet below is how we allocate the memory for writing the input file content to the array in the memory and how we allocate the memory for the output file array.

The line 70 to 73 shows that. Note here is that the data and results variable containing the addresses of the first byte of the two arrays (one for input file content and the other is for the output file). These two pointers and the length of the whole input file array will be passing to the processData function and this function is for the simulation of the logical operation. Note here the length of the results array is $\frac{length}{3}$ and the length of the data array is *length*. The reason of that is for each line in the input file contains 6 characters while for each line in the output file contains 2 characters (one if the result of the logical operation and the other is the new line sign '\n'). The start and end variables are for recording the timing.

```

67 void parseFile(FILE* fp, int length, FILE* output)
68 {
69     char* data;
70     char* results;
71     data = (char*)malloc(length);
72     results = (char*)malloc(length / 3);
73     fread(data, 1, length, fp);
74     clock_t start = clock();
75     processData(data, length, results);
76     clock_t end = clock();
77     fputs(results, output);
78     // printf("Clocks per second == %f ", CLOCKS_PER_SEC);
79     printf("Time used: %f\n", (float)(end-start)/CLOCKS_PER_SEC);
80     free(data);
81     free(results);
82 }
83

```

Figure 1.2: CPU Code Snippet parseFile

The code snippet below shows the processData function. The for loop will iterate every element of the results array for the output file. For each line in the output file there are two characters one is the result of the logical operation and the other is the new line sign '\n' so in each if block we process one line of the input file. We first need to get which logic gate for this line (AND, NAND, OR, NOR, XOR, or XNOR). Each of them is representing by a number and the mapping is shown below given in the lab manual. The fifth element of each line in the input file is the type of logic gate for this operation. We first compare the data[i*3+4] with the number representing each logic gate. The reason we multiply i by 3 is that the length of each input file line is 6 and the length of each output file length is 2 and we are iterate every element of the array for the output file in order to map the input file array's element index with the output file array's element index we need to multiply the input file array's element index by 3.

```

14 void processData(char* data, int length, char* results)
15 {
16     for(int i=0; i<length/3; i++)
17     {
18         // a = data[i];
19         // b = data[i + 2];
20         // opCode = data[i + 4];
21         // in ascii code 1 is 49 and 0 is 48
22         if(data[i*3 + 4] == '0')
23         {
24             int result = (((data[i*3]) - '0') & ((data[i*3 + 2]) - '0'));
25             results[i] = (result + '0');
26             i++;
27             results[i] = '\n';
28         }
29         else if(data[i*3 + 4] == '1')
30         {
31             int result = (((data[i*3]) - '0') | ((data[i*3 + 2]) - '0'));
32             results[i] = (result + '0');
33             i++;
34             results[i] = '\n';
35         }
36         else if(data[i*3 + 4] == '2')
37         {
38             int result = !(((data[i*3]) - '0') & ((data[i*3 + 2]) - '0'));
39             results[i] = (result + '0');
40             i++;
41             results[i] = '\n';
42         }
43         else if(data[i*3 + 4] == '3')
44         {
45             int result = !(((data[i*3]) - '0') | ((data[i*3 + 2]) - '0'));
46             results[i] = (result + '0');
47             i++;
48             results[i] = '\n';
49         }
50         else if(data[i*3 + 4] == '4')
51         {
52             int result = (((data[i*3]) - '0') ^ ((data[i*3 + 2]) - '0'));
53             results[i] = (result + '0');
54             i++;
55             results[i] = '\n';
56         }
57         else if(data[i*3 + 4] == '5')
58         {
59             int result = !(((data[i*3]) - '0') ^ ((data[i*3 + 2]) - '0'));
60             results[i] = (result + '0');
61             i++;
62             results[i] = '\n';
63         }
64     }
65 }

```

Figure 1.3 : CPU Code Snippet Simulation

```

# define AND 0
# define OR 1
# define NAND 2
# define NOR 3
# define XOR 4
# define XNOR 5

```

Figure 1.4 : CPU Code Logical Operations Representation

In figure below we can see the arrangement of the input array and the line in the input file. The last character is for csv file to represent this line is at the end. As shown in the figure, data[i*3+0] and data[i*3+2] are the first and second operands, the data[i*3+1] and data[i*3+3] are the comma for the csv file. The data[i*3+5] is the new line sign.

data[i*3+0]	data[i*3+1]	data[i*3+2]	data[i*3+3]	data[i*3+4]	data[i*3+5]
0	,	1	,	4	'\n'

Figure 1.5: Array Element Mapping With Characters In Each Line Of Input File

The reason in each if block that we need a `i++` is that we need to get the index for putting the new line sign in the output file array. Because C is using ascii code and the '0' and '1' are next to each other ('0' in the ascii code is 48 and '1' is 49) so we simply need to use the value inside the input array to minus the '0' character to cast the character to integer and then do the logical operation on those to make the simulation. From the official C document, we know that how to do these logical operations in C and it is shown in the table below.

Logical Operation	AND	OR	NAND	NOR	XOR	XNOR
Bitwise Operation In C	&		!(op1&op2)	!(op1 op2)	^	!(op1^op2)

Result:

The figure below shows the comparison result between our output files and the solutions given on mycourses. There are no errors which means that our method is right.

```
PS C:\Users\as785\source\repos\ECSE420_Lab1\Compare> ./compare sol_10000.txt sequential_output_10000.txt
Total Errors : 0
PS C:\Users\as785\source\repos\ECSE420_Lab1\Compare> ./compare sol_100000.txt sequential_output_100000.txt
Total Errors : 0
PS C:\Users\as785\source\repos\ECSE420_Lab1\Compare> ./compare sol_1000000.txt sequential_output_1000000.txt
Total Errors : 0
```

Figure 1.6: Results of comparison between explicit outputs and solutions

The figure below shows the time record for the sequential part. Note here we used the built in C library `time.h` and use the function `clock()` inside the library. The `clock()` will give us how many cycles that the CPU clock has run right now so we need to use the number of clock cycles that the CPU has been used after the `processData` function subtract the number of clock cycles that the CPU has been used right before the `processData`. Here we have a convention issue. Since what we got is the number of clock cycle we need to divide it by the constant `CLOCKS_PER_SEC` but this constant is different between macOS (Unix), Linux and Windows. For the figure below it is running on macOS and it is representing microsecond.

```
[(base) Riverflixs-Mac:Lab1 zengmai-river$ ./sequential Resources/input_10000.csv 10000 output10000.txt
Time used: 0.169000 ms
[(base) Riverflixs-Mac:Lab1 zengmai-river$ ./sequential Resources/input_100000.csv 100000 output100000.txt
Time used: 2.381000 ms
[(base) Riverflixs-Mac:Lab1 zengmai-river$ ./sequential Resources/input_1000000.csv 1000000 output1000000.txt
Time used: 22.067999 ms
```

Figure 1.7: Host function times on macOS

2. Parallelize your code using explicit memory allocation in CUDA:

Code:

In this question, we are asked to parallelize our code using explicit memory allocation. As a result, we need to do memory allocation in both CPU and GPU. We first use `malloc` to allocate memory in CPU and `cudaMalloc` to do that in GPU. Then we use

cudaMemcpy to copy memory from CPU to GPU (cudaMemcpyHostToDevice) in order to ensure that their memories are equal. A timer is started before this cudaMemcpy function and stops immediately after memory copy is done. This time period is the explicit data migration time. This part of code is shown below.

```
void parallel_explicit(FILE* fp_in, int length, FILE* fp_out) {
    // input has length 'length' and output has length 'length/3'
    // output file has only 2 elements in one line (a number and a '\n') while input file has 6 (
    char* data, * d_data, * results, * d_results;
    // timer_kernel records time for kernel function
    // timer_migration records explicit data migration time (copy data from host to device)
    GpuTimer timer_kernel, timer_migration;
    data = (char*)malloc(length);
    results = (char*)malloc(length/3);
    cudaMalloc(&d_data, length);
    cudaMalloc(&d_results, length/3);
    fread(data, 1, length, fp_in);
    timer_migration.Start();
    cudaMemcpy(d_data, data, length, cudaMemcpyHostToDevice);
    cudaMemcpy(d_results, results, length/3, cudaMemcpyHostToDevice);
    timer_migration.Stop();
}
```

Figure 2.1: Explicit Code Snippet

Since the input file has a great number of lines, which is also number of logic gates, and the maximum number of threads per block is 1024, we use the method “parallel blocks and parallel threads” in the kernel function. This is implemented as $i = \text{threadIdx.x} + \text{blockIdx.x} * \text{blockDim.x}$. Since our program only needs to launch one thread per logic gate, which is also one thread per line of the input file, this i should be less than the total number of lines in the file to ensure the prerequisite. One line of an input file has 6 characters, the first and third numbers are inputs, the fifth number represents its corresponding gate, the second and fourth numbers are commas while the sixth character is a newline character. Consequently, we first check the fifth number in each line, which should be $[i * 6 + 4]$, after that we check the inputs and conclude a result according to the truth table, finally we store it into the output at position $[2 * i]$ with a newline character at $[2 * i + 1]$. This part of code is shown below.

```

__global__ void classify(char* d_data, int SIZE, char* d_results) {
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    // since only needs to launch one thread per logic gate, SIZE should equal to the number of rows in the file
    if (i < SIZE) {
        if (d_data[i * 6 + 4] == '0')
        {
            int result = (((d_data[i * 6]) - '0') & ((d_data[i * 6 + 2]) - '0'));
            d_results[2 * i] = (result + '0');
            d_results[2 * i + 1] = '\n';
        }
        else if (d_data[i * 6 + 4] == '1')
        {
            int result = (((d_data[i * 6]) - '0') | ((d_data[i * 6 + 2]) - '0'));
            d_results[2 * i] = (result + '0');
            d_results[2 * i + 1] = '\n';
        }
        else if (d_data[i * 6 + 4] == '2')
        {
            int result = !(((d_data[i * 6]) - '0') & ((d_data[i * 6 + 2]) - '0'));
            d_results[2 * i] = (result + '0');
            d_results[2 * i + 1] = '\n';
        }
        else if (d_data[i * 6 + 4] == '3')
        {
            int result = !(((d_data[i * 6]) - '0') | ((d_data[i * 6 + 2]) - '0'));
            d_results[2 * i] = (result + '0');
            d_results[2 * i + 1] = '\n';
        }
        else if (d_data[i * 6 + 4] == '4')
        {
            int result = (((d_data[i * 6]) - '0') ^ ((d_data[i * 6 + 2]) - '0'));
            d_results[2 * i] = (result + '0');
            d_results[2 * i + 1] = '\n';
        }
        else if (d_data[i * 6 + 4] == '5')
        {
            int result = !(((d_data[i * 6]) - '0') ^ ((d_data[i * 6 + 2]) - '0'));
            d_results[2 * i] = (result + '0');
            d_results[2 * i + 1] = '\n';
        }
    }
}

```

Figure 2.2: Explicit Code Snippet Simulation

Before we call this kernel function on GPU, we need to pass in two parameters, totalBlocks and maxThreadNum. Our idea is that we use the number of lines to divide the maxThreadNum, whose default value is set to 1024, we continue decrease this value by one until this division returns a remainder 0. In this situation, the total number of lines is equally distributed to all the blocks we used. A timer is also set before and after we call the kernel function to record its execution time. After we finish all the execution on GPU, cudaMemcpy is used again but this time it is from GPU to CPU (cudaMemcpyDeviceToHost), its aim is still to ensure the memories on GPU and CPU are equal. At last, we use cudaFree() to free memory on GPU and free() to do that on CPU. This part of code is shown below, it is in the same class and right after the code shown first in this question.

```

int maxThreadNum = 1024;
// distribute the total threads equally in blocks
while(1)
{
    if(length/6 % maxThreadNum != 0)
    {
        maxThreadNum--;
    }
    else
    {
        break;
    }
}
int totalBlocks = length / 6 / maxThreadNum;
timer_kernel.Start();
classify <<<totalBlocks, maxThreadNum>>> (d_data, length/6, d_results);
timer_kernel.Stop();
cudaMemcpy(data, d_data, length, cudaMemcpyDeviceToHost);
cudaMemcpy(results, d_results, length/3, cudaMemcpyDeviceToHost);
printf("Time for kernel functions: %f ms\n", timer_kernel.Elapsed());
printf("Time for explicit data migration: %f ms\n", timer_migration.Elapsed());

fputs(results, fp_out);

cudaFree(d_data);
cudaFree(d_results);
free(data);
free(results);
}

```

Figure 2.3: Explicit Code Snippet Main

Result:

The times for three input files are listed below:

```

For input file ./input_10000.csv
Time for kernel functions: 0.016384 ms
Time for explicit data migration: 0.193216 ms

```

Figure 2.4: Necessary times for input file input_10000.csv

```

For input file ./input_100000.csv
Time for kernel functions: 0.057856 ms
Time for explicit data migration: 0.255104 ms

```

Figure 2.5: Necessary times for input file input_100000.csv

```

For input file ./input_1000000.csv
Time for kernel functions: 0.448512 ms
Time for explicit data migration: 1.378400 ms

```

Figure 2.6: Necessary times for input file input_1000000.csv

The compare results between three output files and solutions are listed below:

```

PS C:\Users\8785\source\repos\ECSE420_Lab1\Compare> ./compare sol_10000.txt explicit_output_10000.txt
Total Errors : 0
PS C:\Users\8785\source\repos\ECSE420_Lab1\Compare> ./compare sol_100000.txt explicit_output_100000.txt
Total Errors : 0
PS C:\Users\8785\source\repos\ECSE420_Lab1\Compare> ./compare sol_1000000.txt explicit_output_1000000.txt
Total Errors : 0

```

Figure 2.7: Results of comparison between explicit outputs and solutions

This means that our function is correct.

3. Parallelize your code using unified memory allocation in CUDA:

Code:

In this question, we are asked to parallelize the code using unified memory allocation. In a unified memory, if we understand the concept in a simple way, that is both CPU and GPU share the same memory space. As a result, we do not need to use two different methods to malloc space, keep updating them by copying memory and use two different ways to free them at last, instead we use the method `cudaMallocManaged` to malloc the space once and `cudaDeviceSynchronize()` to force the program to ensure the streams' kernels are complete before continuing, we only use `cudaFree()` once at last. Apart from this change, the remaining logics are the same as question 2. There is no data migration time because this is not needed in unified memory allocation. The part of code that is changed is shown below.

```
void parallel_unified(FILE* fp_in, int length, FILE* fp_out) {
    // input has length 'length' and output has length 'length/3'
    // output file has only 2 elements in one line (a number and a '\n') while input file has 6 (listed in main)
    char* data, * results;
    // timer_kernel records time for kernel function
    GpuTimer timer_kernel;
    // Unified memory allocation methods
    cudaMallocManaged(&data, length);
    cudaMallocManaged(&results, length / 3);
    fread(data, 1, length, fp_in);
    int maxThreadNum = 1024;
    // distribute the total threads equally in blocks
    while (1)
    {
        if (length / 6 % maxThreadNum != 0)
        {
            maxThreadNum--;
        }
        else
        {
            break;
        }
    }
    int totalBlocks = length / 6 / maxThreadNum;
    timer_kernel.Start();
    classify <<<totalBlocks, maxThreadNum >>> (data, length / 6, results);
    timer_kernel.Stop();
    cudaDeviceSynchronize();
    printf("Time for kernel functions: %f ms\n", timer_kernel.Elapsed()); // Convert unit from ms to s

    fputs(results, fp_out);

    cudaFree(data);
    cudaFree(results);
}
```

Figure 3.1: Unified Code Snippet

Result:

The times for three input files are listed below:

```
For input file ./input_10000.csv
Time for kernel functions: 0.017184 ms
```

Figure 3.2: Kernel function time for input file input_10000.csv

```
For input file ./input_100000.csv
Time for kernel functions: 0.059744 ms
```

Figure 3.3: Kernel function time for input file input_100000.csv

```
For input file ./input_1000000.csv
Time for kernel functions: 0.449312 ms
```

Figure 3.4: Kernel function time for input file input_1000000.csv

The compare results between three output files and solutions are listed below:

```
PS C:\Users\A8785\source\repos\ECSE420_Lab1\Compare> ./compare sol_10000.txt unified_output_10000.txt
Total Errors : 0
PS C:\Users\A8785\source\repos\ECSE420_Lab1\Compare> ./compare sol_100000.txt unified_output_100000.txt
Total Errors : 0
PS C:\Users\A8785\source\repos\ECSE420_Lab1\Compare> ./compare sol_1000000.txt unified_output_1000000.txt
Total Errors : 0
```

Figure 3.5: Results of comparison between unified outputs and solutions

This means that our function is correct. For all the kernel functions' execution times, they are a little bit larger (around 0.1 ms to 0.2 ms) in unified memory compared to those in explicit memory.

As we can see in the timing records, the unified memory is taking slightly longer than explicit parallel method running the kernel function. More details will be talked in the fifth section.

4. Parallel code using unified memory allocation with data prefetching in CUDA:

Since the CUDA windows version is not supporting the prefetching function we are going to use colab for this part.

The code implementation is simple. We just need to add the code snippet right before we call the kernel function like the Figure 4.1 below.

```
// Prefetch the data to the GPU
int device = -1;
cudaGetDevice(&device);
cudaMemPrefetchAsync(data, length, device, NULL);
cudaMemPrefetchAsync(results, length/3, device, NULL);

timer_kernel.Start();
classify <<<totalBlocks, maxThreadNum >>> (data, length / 6, results);
timer_kernel.Stop();
cudaDeviceSynchronize();
printf("Time for kernel functions: %f ms\n", timer_kernel.Elapsed());
```

Figure 4.1: Code Snippet For Unified Memory With Prefetching

In this way we can first prefetching the memory in Unified memory for GPU to avoid the page fault for GPU memory which will be talked in the latter section.

The result for prefetching is shown below. Note here that since it is running on Google Colab so the time may not be comparable with the previous part. More details will be talked in the latter section.

```
[24] !nvcc parallel_prefetch.cu -o parallel_prefetch
!./parallel_prefetch input_10000.csv 10000 output_parallel_prefetch_10000.txt
!./parallel_prefetch input_100000.csv 100000 output_parallel_prefetch_100000.txt
!./parallel_prefetch input_1000000.csv 1000000 output_parallel_prefetch_1000000.txt

For input file input_10000.csv
Time for kernel functions: 0.011680 ms
For input file input_100000.csv
Time for kernel functions: 0.022592 ms
For input file input_1000000.csv
Time for kernel functions: 0.124448 ms
```

Figure 4.2: Time Recording For Unified Memory With Prefetching

Comparison For Unified Memory With Prefetching

```
!gcc compareResults.c -o compareResults
!./compareResults output_parallel_prefetch_10000.txt sol_10000.txt
!./compareResults output_parallel_prefetch_100000.txt sol_100000.txt
!./compareResults output_parallel_prefetch_1000000.txt sol_1000000.txt

Total Errors : 0      Total Errors : 0      Total Errors : 0
```

Figure 4.3: Comparison Between Output Files and Solutions

5. Discuss the results for each architecture:

We did the experiments using the CPU code with sequential method, CUDA code with explicit memory allocation method (Ordinary CUDA code), and CUDA code with unified memory.

The results are shown below:

For explicit parallel:

```
For input file ./input_10000.csv
Time for kernel functions: 0.016384 ms
Time for explicit data migration: 0.193216 ms
For input file ./input_100000.csv
Time for kernel functions: 0.057856 ms
Time for explicit data migration: 0.255104 ms
For input file ./input_1000000.csv
Time for kernel functions: 0.448512 ms
Time for explicit data migration: 1.378400 ms
```

Figure 5.1: Time Recording For Executing csv Files For Explicit Parallel

For unified parallel:

```
For input file ./input_10000.csv
Time for kernel functions: 0.017184 ms
For input file ./input_100000.csv
Time for kernel functions: 0.059744 ms
For input file ./input_1000000.csv
Time for kernel functions: 0.449312 ms
```

Figure 5.2: Time Recording For Executing csv Files For Unified Parallel

For sequential part:

```
(base) Riverflixs-Mac:Lab1 zengmai-river$ ./sequential Resources/input_10000.csv 10000 output10000.txt
Time used: 0.169000 ms
(base) Riverflixs-Mac:Lab1 zengmai-river$ ./sequential Resources/input_100000.csv 100000 output100000.txt
Time used: 2.381000 ms
(base) Riverflixs-Mac:Lab1 zengmai-river$ ./sequential Resources/input_1000000.csv 1000000 output1000000.txt
Time used: 22.067999 ms
```

Figure 5.3: Time Recording For Executing csv Files For CPU Code

As we can see from the timing records. The code running only on CPU is much slower than the CUDA code with either the methods. Comparing the two CUDA codes, with the unified parallel method the time for running the kernel functions is a little bit longer but almost the same as the time running the kernel functions with explicit parallel method for both 3 files but note here the parallel explicit method need the data migration time.

The explicit memory architecture can be represented like the figure below:

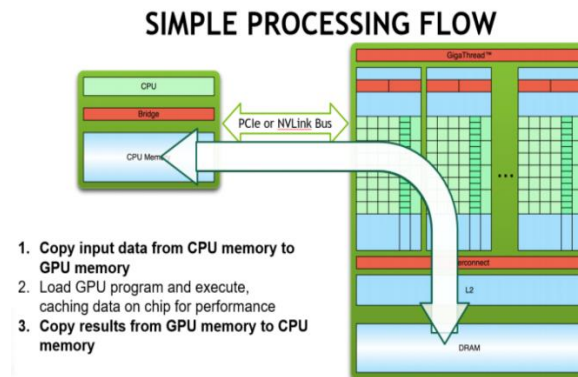


Figure 5.4: Explicit memory architecture

The data first need to copy from CPU memory to GPU memory and that is why data migration taking a significant amount of time. After the data migration, it will load the GPU program and then execute.

The unified memory architecture can be represented like the figure below:

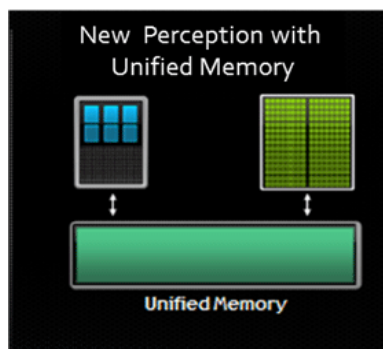


Figure 5.5: Unified memory architecture

As shown in the Figure 5.5, unified Memory is a single memory address space accessible from any processor in a system. This hardware/software technology allows applications to allocate data that can be read or written from code running on either CPUs or GPUs. When code running on a CPU or GPU accesses data allocated this way (often called CUDA managed data), the CUDA system software and/or the hardware takes care of migrating memory pages to the memory of the accessing processor. [1]

CUDA 6 adds one extra layer of convenience to the CPU/GPU memory management task with the introduction of unified or managed memory. Data is now stored and migrated in a user-transparent fashion that enables, under circumstances spelled out shortly, data access/transfer at latencies and bandwidths of the host and of the device, for host-side and device-side memory operations, respectively. Moreover, the use of the `cudaHostAlloc` and `cudaMemcpy` combination is no longer a requirement, which allows for a cleaner and more natural programming style.

This means that after CUDA6 the developer does not need to concern the `cudaMemcpy` and `cudaMalloc` instead they can just use `cudaMallocManaged` function. This eliminate the need for explicit copy and it still allows explicit hand tuning.

The reason that the unified memory is a little bit slower than the explicit is that the memory is unified but it still need to communicate between the GPU memory with the unified memory and that will take some time but with the unified memory, it does not need the data migration time which is actually a significant amount of time.

For better understand the difference between explicit parallel, unified parallel and prefetch with unified memory. We tested our code on the Google Colab and run the command 'nvprof' to see the whole running profile.

For explicit parallel without using unified memory:

```
nvprof ./parallel_explicit input_10000.csv 10000 output_parallel_explicit_10000.txt
```

For input file input_10000.csv
 ==509== NVPROF is profiling process 509, command: ./parallel_explicit input_10000.csv 10000 output_parallel_explicit_10000.txt
 Time for kernel functions: 0.027392 ms
 Time for explicit data migration: 0.112224 ms
 ==509== Profiling application: ./parallel_explicit input_10000.csv 10000 output_parallel_explicit_10000.txt
 ==509== Profiling result:

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	47.38%	11.840us	2	5.9200us	3.5840us	8.2560us	[CUDA memcpy HtoD]
	31.63%	7.9040us	2	3.9520us	2.4000us	5.5040us	[CUDA memcpy DtoH]
	21.00%	5.2480us	1	5.2480us	5.2480us	5.2480us	classify(char*, int, char*)
API calls:	99.42%	167.03ms	6	27.838ms	565ns	167.02ms	cudaEventCreate
	0.22%	368.95us	1	368.95us	368.95us	368.95us	cudaDeviceTotalMem
	0.10%	170.71us	2	85.353us	6.7840us	163.92us	cudaMalloc
	0.08%	137.51us	97	1.4170us	134ns	52.172us	cudaDeviceGetAttribute
	0.06%	106.27us	4	26.568us	18.283us	32.950us	cudaMemcpy
	0.06%	102.40us	2	51.198us	15.017us	87.379us	cudaFree
	0.02%	25.548us	1	25.548us	25.548us	25.548us	cudaLaunchKernel
	0.01%	22.767us	1	22.767us	22.767us	22.767us	cudaDeviceGetName
	0.01%	14.466us	6	2.4110us	1.2370us	6.1700us	cudaEventRecord
	0.00%	8.2990us	3	2.7660us	1.7550us	3.3620us	cudaEventElapsedTime
	0.00%	7.3050us	3	2.4350us	2.1230us	2.6880us	cudaEventSynchronize
	0.00%	4.7900us	6	798ns	427ns	1.8140us	cudaEventDestroy
	0.00%	3.8250us	1	3.8250us	3.8250us	3.8250us	cudaDeviceGetPCIBusId
	0.00%	2.1390us	3	713ns	140ns	1.2460us	cudaDeviceGetCount
	0.00%	1.5590us	2	779ns	393ns	1.1660us	cudaDeviceGet
	0.00%	250ns	1	250ns	250ns	250ns	cudaDeviceGetUuid

Figure 5.6: Explicit 10000

```

nvprof ./parallel_explicit input_100000.csv 100000 output_parallel_explicit_100000.txt

For input file input_100000.csv
==524== NVPROF is profiling process 524, command: ./parallel_explicit input_100000.csv 100000 output_parallel_explicit_100000.txt
Time for kernel functions: 0.021568 ms
Time for explicit data migration: 0.559360 ms
==524== Profiling application: ./parallel_explicit input_100000.csv 100000 output_parallel_explicit_100000.txt
==524== Profiling result:
Type Time(%) Time Calls Avg Min Max Name
GPU activities: 48.13% 70.655us 2 35.327us 18.688us 51.967us [CUDA memcpy HtoD]
45.79% 67.231us 2 33.615us 16.384us 50.847us [CUDA memcpy DtoH]
6.08% 8.9280us 1 8.9280us 8.9280us 8.9280us classify(char*, int, char*)
API calls: 99.12% 156.56ms 6 26.094ms 952ns 156.56ms cudaEventCreate
0.34% 544.91us 4 136.23us 104.78us 184.16us cudaMemcpy
0.25% 392.76us 1 392.76us 392.76us 392.76us cuDeviceTotalMem
0.08% 130.25us 2 65.122us 7.1480us 123.10us cudaMalloc
0.08% 130.00us 97 1.3400us 136ns 48.297us cuDeviceGetAttribute
0.07% 105.88us 2 52.942us 16.223us 89.661us cudaFree
0.02% 27.656us 1 27.656us 27.656us 27.656us cudaLaunchKernel
0.01% 16.805us 6 2.8000us 1.2240us 8.0460us cudaEventRecord
0.01% 14.874us 1 14.874us 14.874us 14.874us cuDeviceGetName
0.01% 8.8720us 3 2.9570us 1.8220us 3.6410us cudaEventElapsedTime
0.00% 7.7440us 3 2.5810us 2.1890us 2.8550us cudaEventSynchronize
0.00% 6.1340us 6 1.0220us 453ns 2.9440us cudaEventDestroy
0.00% 3.2270us 1 3.2270us 3.2270us 3.2270us cuDeviceGetPCIBusId
0.00% 2.3370us 3 779ns 167ns 1.6030us cuDeviceGetCount
0.00% 1.1160us 2 558ns 265ns 851ns cuDeviceGet
0.00% 266ns 1 266ns 266ns 266ns cuDeviceGetUuid

```

Figure 5.7: Explicit 100000

```

nvprof ./parallel_explicit input_1000000.csv 1000000 output_parallel_explicit_1000000.txt

For input file input_1000000.csv
==537== NVPROF is profiling process 537, command: ./parallel_explicit input_1000000.csv 1000000 output_parallel_explicit_1000000.txt
Time for kernel functions: 0.073632 ms
Time for explicit data migration: 4.961696 ms
==537== Profiling application: ./parallel_explicit input_1000000.csv 1000000 output_parallel_explicit_1000000.txt
==537== Profiling result:
Type Time(%) Time Calls Avg Min Max Name
GPU activities: 61.31% 1.5204ms 2 760.20us 398.27us 1.1221ms [CUDA memcpy HtoD]
36.06% 894.13us 2 447.06us 155.17us 738.96us [CUDA memcpy DtoH]
2.63% 65.215us 1 65.215us 65.215us 65.215us classify(char*, int, char*)
API calls: 96.38% 166.19ms 6 27.698ms 542ns 166.18ms cudaEventCreate
2.89% 4.9862ms 4 1.2465ms 755.75us 1.8855ms cudaMemcpy
0.24% 407.62us 1 407.62us 407.62us 407.62us cuDeviceTotalMem
0.22% 378.29us 2 189.15us 115.91us 262.39us cudaFree
0.14% 233.65us 2 116.82us 108.19us 125.46us cudaMalloc
0.08% 133.21us 97 1.3730us 142ns 49.563us cuDeviceGetAttribute
0.02% 30.196us 1 30.196us 30.196us 30.196us cudaLaunchKernel
0.02% 27.501us 6 4.5830us 1.2220us 12.873us cudaEventRecord
0.01% 17.041us 1 17.041us 17.041us 17.041us cuDeviceGetName
0.00% 8.2150us 3 2.7380us 1.6510us 3.3190us cudaEventElapsedTime
0.00% 8.1830us 3 2.7270us 2.2850us 3.0850us cudaEventSynchronize
0.00% 6.5630us 6 1.0930us 464ns 3.3840us cudaEventDestroy
0.00% 2.9120us 1 2.9120us 2.9120us 2.9120us cuDeviceGetPCIBusId
0.00% 1.9810us 3 660ns 183ns 1.1260us cuDeviceGetCount
0.00% 1.2200us 2 610ns 360ns 860ns cuDeviceGet
0.00% 303ns 1 303ns 303ns 303ns cuDeviceGetUuid

```

Figure 5.8: Explicit 1000000

For Unified memory without prefetching:

```

nvprof ./parallel_unified input_10000.csv 10000 output_parallel_unified_10000.txt

For input file input_10000.csv
==790== NVPROF is profiling process 790, command: ./parallel_unified input_10000.csv 10000 output_parallel_unified_10000.txt
Time for kernel functions: 0.605248 ms
==790== Profiling application: ./parallel_unified input_10000.csv 10000 output_parallel_unified_10000.txt
==790== Profiling result:
Type Time(%) Time Calls Avg Min Max Name
GPU activities: 100.00% 556.76us 1 556.76us 556.76us 556.76us classify(char*, int, char*)
API calls: 88.08% 160.93ms 2 80.466ms 948ns 160.93ms cudaEventCreate
11.17% 20.407ms 2 10.204ms 34.904us 20.373ms cudaMallocManaged
0.31% 565.49us 1 565.49us 565.49us 565.49us cudaDeviceSynchronize
0.24% 430.73us 1 430.73us 430.73us 430.73us cuDeviceTotalMem
0.08% 146.94us 97 1.5140us 135ns 47.650us cuDeviceGetAttribute
0.07% 134.20us 2 67.098us 28.889us 105.31us cudaFree
0.03% 45.986us 1 45.986us 45.986us 45.986us cudaLaunchKernel
0.01% 19.309us 2 9.6540us 3.0040us 16.305us cudaEventRecord
0.01% 13.381us 1 13.381us 13.381us 13.381us cuDeviceGetName
0.00% 3.5650us 1 3.5650us 3.5650us 3.5650us cudaEventSynchronize
0.00% 3.3330us 2 1.6660us 610ns 2.7230us cudaEventDestroy
0.00% 3.1900us 1 3.1900us 3.1900us 3.1900us cuDeviceGetPCIBusId
0.00% 3.0890us 1 3.0890us 3.0890us 3.0890us cudaEventElapsedTime
0.00% 1.7780us 3 592ns 163ns 1.1050us cuDeviceGetCount
0.00% 1.1700us 2 585ns 240ns 930ns cuDeviceGet
0.00% 274ns 1 274ns 274ns 274ns cuDeviceGetUuid

==790== Unified Memory profiling result:
Device "Tesla P100-PCIE-16GB (0)"
Count Avg Size Min Size Max Size Total Size Total Time Name
2 32.000KB 4.0000KB 60.000KB 64.00000KB 11.16800us Host To Device
5 25.600KB 4.0000KB 60.000KB 128.0000KB 22.78400us Device To Host
2 - - - 549.0880us Gpu page fault groups
Total CPU Page faults: 3

```

Figure 5.9: Unified without prefetching 10000


```

nvprof ./parallel_unified input_100000.csv 100000 output_parallel_unified_100000.txt

For input file input_100000.csv
==777== NVPROF is profiling process 777, command: ./parallel_unified input_100000.csv 100000 output_parallel_unified_100000.txt
Time for kernel functions: 0.676928 ms
==777== Profiling application: ./parallel_unified input_100000.csv 100000 output_parallel_unified_100000.txt
==777== Profiling result:
   Type  Time(%)      Time   Calls    Avg      Min      Max   Name
GPU activities: 100.00%  637.56us    1  637.56us  637.56us  637.56us  classify(char*, int, char*)
API calls:      88.38%  165.26ms    2   82.630ms  1.0250us  165.26ms  cudaEventCreate
              10.85%  20.297ms    2   10.148ms  25.071us  20.272ms  cudaMallocManaged
              0.30%   568.94us    1   568.94us  568.94us  568.94us  cudaDeviceSynchronize
              0.24%   445.93us    1   445.93us  445.93us  445.93us  cuDeviceTotalMem
              0.11%   200.83us    2   100.42us  69.844us  130.99us  cudaFree
              0.07%   132.86us    97   1.3690us   148ns   48.198us  cuDeviceGetAttribute
              0.02%   33.550us    1   33.550us  33.550us  33.550us  cudaLaunchKernel
              0.01%   16.952us    2    8.4760us  4.6210us  12.331us  cudaEventRecord
              0.01%   14.174us    1   14.174us  14.174us  14.174us  cuDeviceGetName
              0.00%   3.9760us    1   3.9760us  3.9760us  3.9760us  cudaEventSynchronize
              0.00%   3.6880us    2   1.8440us   651ns   3.0370us  cudaEventDestroy
              0.00%   3.5430us    1   3.5430us  3.5430us  3.5430us  cudaEventElapsedTime
              0.00%   2.6510us    1   2.6510us  2.6510us  2.6510us  cuDeviceGetPCIBusId
              0.00%   1.8170us    3     605ns    197ns   1.2240us  cuDeviceGetCount
              0.00%   1.0930us    2     546ns    244ns    849ns  cuDeviceGet
              0.00%    280ns    1     280ns    280ns    280ns  cuDeviceGetUuid

==777== Unified Memory profiling result:
Device "Tesla P100-PCIE-16GB (0)"
Count Avg Size Min Size Max Size Total Size Total Time Name
  38  26.947KB  4.0000KB 244.00KB 1.000000MB 144.0000us Host To Device
  10  51.199KB  4.0000KB 252.00KB 512.0000KB 49.63200us Device To Host
   2  - - - - 620.3840us Cpu page fault groups
Total CPU Page faults: 9

```

Figure 5.10: Unified without prefetching 100000

```

nvprof ./parallel_unified input_1000000.csv 1000000 output_parallel_unified_1000000.txt

For input file input_1000000.csv
==801== NVPROF is profiling process 801, command: ./parallel_unified input_1000000.csv 1000000 output_parallel_unified_1000000.txt
Time for kernel functions: 2.741632 ms
==801== Profiling application: ./parallel_unified input_1000000.csv 1000000 output_parallel_unified_1000000.txt
==801== Profiling result:
   Type  Time(%)      Time   Calls    Avg      Min      Max   Name
GPU activities: 100.00%  2.6897ms    1  2.6897ms  2.6897ms  2.6897ms  classify(char*, int, char*)
API calls:      87.52%  170.18ms    2   85.089ms   977ns  170.18ms  cudaEventCreate
              10.52%  20.459ms    2   10.230ms  43.812us  20.416ms  cudaMallocManaged
              1.39%   2.7029ms    1   2.7029ms  2.7029ms  2.7029ms  cudaDeviceSynchronize
              0.22%   434.79us    2   217.39us  121.24us  313.55us  cudaFree
              0.21%   415.24us    1   415.24us  415.24us  415.24us  cuDeviceTotalMem
              0.07%   142.17us    97   1.4650us   158ns   54.434us  cuDeviceGetAttribute
              0.03%   50.315us    1   50.315us  50.315us  50.315us  cudaLaunchKernel
              0.01%   17.454us    1   17.454us  17.454us  17.454us  cuDeviceGetName
              0.01%   15.674us    2    7.8370us  3.1340us  12.540us  cudaEventRecord
              0.00%   3.9970us    1   3.9970us  3.9970us  3.9970us  cudaEventSynchronize
              0.00%   3.6670us    2    1.8330us   642ns   3.0250us  cudaEventDestroy
              0.00%   3.2710us    1   3.2710us  3.2710us  3.2710us  cuDeviceGetPCIBusId
              0.00%   3.1350us    1   3.1350us  3.1350us  3.1350us  cudaEventElapsedTime
              0.00%   2.2370us    3     745ns    200ns   1.4080us  cuDeviceGetCount
              0.00%   1.2720us    2     636ns    274ns    998ns  cuDeviceGet
              0.00%    317ns    1     317ns    317ns    317ns  cuDeviceGetUuid

==801== Unified Memory profiling result:
Device "Tesla P100-PCIE-16GB (0)"
Count Avg Size Min Size Max Size Total Size Total Time Name
  239  24.519KB  4.0000KB  724.00KB 5.722656MB 813.1840us Host To Device
   12  170.67KB  4.0000KB  0.9961MB 2.000000MB 173.6640us Device To Host
    16  - - - - 2.901920ms Cpu page fault groups
Total CPU Page faults: 24

```

Figure 5.11: Unified without prefetching 1000000

For Unified Memory with Prefetching:

```

▶ !nvprof ./parallel_prefetch input_10000.csv 10000 output_parallel_prefetch_10000.txt

┌ For input file input_10000.csv
==361== NVPROF is profiling process 361, command: ./parallel_prefetch input_10000.csv 10000 output_parallel_prefetch_10000.txt
Time for kernel functions: 0.042752 ms
==361== Profiling application: ./parallel_prefetch input_10000.csv 10000 output_parallel_prefetch_10000.txt
==361== Profiling result:

```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	100.00%	6.6870us	1	6.6870us	6.6870us	6.6870us	classify(char*, int, char*)
API calls:	88.58%	167.88ms	2	83.938ms	960ns	167.88ms	cudaEventCreate
	10.91%	20.669ms	2	10.335ms	34.143us	20.635ms	cudaMallocManaged
	0.19%	362.81us	1	362.81us	362.81us	362.81us	cudaDeviceTotalMem
	0.13%	252.41us	2	126.20us	21.344us	231.06us	cudaMemPrefetchAsync
	0.07%	138.20us	97	1.4240us	133ns	58.893us	cudaDeviceGetAttribute
	0.06%	109.14us	2	54.572us	29.121us	80.023us	cudaFree
	0.02%	35.367us	1	35.367us	35.367us	35.367us	cudaLaunchKernel
	0.01%	27.919us	1	27.919us	27.919us	27.919us	cudaDeviceGetName
	0.01%	14.578us	2	7.2890us	2.8190us	11.759us	cudaEventRecord
	0.01%	10.446us	1	10.446us	10.446us	10.446us	cudaDeviceSynchronize
	0.00%	3.2220us	1	3.2220us	3.2220us	3.2220us	cudaDeviceGetPCIBusId
	0.00%	3.1670us	2	1.5830us	740ns	2.4270us	cudaEventDestroy
	0.00%	2.5730us	1	2.5730us	2.5730us	2.5730us	cudaGetDevice
	0.00%	2.0400us	1	2.0400us	2.0400us	2.0400us	cudaEventSynchronize
	0.00%	1.7500us	3	583ns	139ns	1.1810us	cudaDeviceGetCount
	0.00%	1.7030us	1	1.7030us	1.7030us	1.7030us	cudaEventElapsedTime
	0.00%	1.0620us	2	531ns	305ns	757ns	cudaDeviceGet
	0.00%	272ns	1	272ns	272ns	272ns	cudaDeviceGetUuid

```

==361== Unified Memory profiling result:
Device "Tesla T4 (0)"
Count Avg Size Min Size Max Size Total Size Total Time Name
2 32.000KB 4.0000KB 60.000KB 64.0000KB 12.76800us Host To Device
5 16.000KB 4.0000KB 56.000KB 80.0000KB 13.66400us Device To Host
Total CPU Page faults: 3

```

Figure 5.12: Unified without prefetching 10000

```

▶ !nvprof ./parallel_prefetch input_100000.csv 100000 output_parallel_prefetch_100000.txt

┌ For input file input_100000.csv
==372== NVPROF is profiling process 372, command: ./parallel_prefetch input_100000.csv 100000 output_parallel_prefetch_100000.txt
Time for kernel functions: 0.022816 ms
==372== Profiling application: ./parallel_prefetch input_100000.csv 100000 output_parallel_prefetch_100000.txt
==372== Profiling result:

```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	100.00%	17.151us	1	17.151us	17.151us	17.151us	classify(char*, int, char*)
API calls:	88.83%	172.28ms	2	86.140ms	1.0980us	172.28ms	cudaEventCreate
	10.54%	20.448ms	2	10.224ms	24.223us	20.424ms	cudaMallocManaged
	0.21%	401.05us	2	200.52us	151.72us	249.33us	cudaMemPrefetchAsync
	0.18%	339.46us	1	339.46us	339.46us	339.46us	cudaDeviceTotalMem
	0.08%	146.59us	2	73.296us	33.429us	113.16us	cudaFree
	0.08%	146.48us	97	1.5100us	151ns	62.122us	cudaDeviceGetAttribute
	0.04%	71.964us	1	71.964us	71.964us	71.964us	cudaDeviceSynchronize
	0.02%	44.569us	1	44.569us	44.569us	44.569us	cudaLaunchKernel
	0.01%	27.757us	1	27.757us	27.757us	27.757us	cudaDeviceGetName
	0.01%	18.041us	2	9.0200us	6.1710us	11.870us	cudaEventRecord
	0.00%	4.3550us	1	4.3550us	4.3550us	4.3550us	cudaDeviceGetPCIBusId
	0.00%	3.9640us	2	1.9820us	905ns	3.0590us	cudaEventDestroy
	0.00%	3.8310us	1	3.8310us	3.8310us	3.8310us	cudaGetDevice
	0.00%	2.7900us	1	2.7900us	2.7900us	2.7900us	cudaEventSynchronize
	0.00%	2.1410us	1	2.1410us	2.1410us	2.1410us	cudaEventElapsedTime
	0.00%	2.0810us	3	693ns	175ns	1.3140us	cudaDeviceGetCount
	0.00%	1.3430us	2	671ns	268ns	1.0750us	cudaDeviceGet
	0.00%	266ns	1	266ns	266ns	266ns	cudaDeviceGetUuid

```

==372== Unified Memory profiling result:
Device "Tesla T4 (0)"
Count Avg Size Min Size Max Size Total Size Total Time Name
2 392.00KB 196.00KB 588.00KB 784.0000KB 73.95200us Host To Device
3 90.666KB 4.0000KB 196.00KB 272.0000KB 25.63200us Device To Host
Total CPU Page faults: 6

```

Figure 5.13: Unified without prefetching 100000

```

▶ !nvprof ./parallel_prefetch input_1000000.csv 1000000 output_parallel_prefetch_1000000.txt

┌ For input file input_1000000.csv
==384== NVPROF is profiling process 384, command: ./parallel_prefetch input_1000000.csv 1000000 output_parallel_prefetch_1000000.txt
Time for kernel functions: 0.170720 ms
==384== Profiling application: ./parallel_prefetch input_1000000.csv 1000000 output_parallel_prefetch_1000000.txt
==384== Profiling result:

```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	100.00%	129.43us	1	129.43us	129.43us	129.43us	classify(char*, int, char*)
API calls:	88.01%	166.18ms	2	83.092ms	962ns	166.18ms	cudaEventCreate
	10.79%	20.369ms	2	10.184ms	52.491us	20.316ms	cudaMallocManaged
	0.45%	857.49us	2	428.74us	25.128us	832.36us	cudaMemPrefetchAsync
	0.33%	621.66us	2	310.83us	166.56us	455.11us	cudaFree
	0.19%	361.84us	1	361.84us	361.84us	361.84us	cudaDeviceTotalMem
	0.09%	171.02us	97	1.7630us	141ns	84.938us	cudaDeviceGetAttribute
	0.07%	131.46us	1	131.46us	131.46us	131.46us	cudaDeviceSynchronize
	0.02%	46.970us	1	46.970us	46.970us	46.970us	cudaDeviceGetName
	0.02%	36.938us	1	36.938us	36.938us	36.938us	cudaLaunchKernel
	0.01%	13.962us	2	6.9810us	4.1210us	9.8410us	cudaEventRecord
	0.00%	4.5980us	1	4.5980us	4.5980us	4.5980us	cudaGetDevice
	0.00%	3.9000us	2	1.9500us	619ns	3.2810us	cudaEventDestroy
	0.00%	3.6390us	1	3.6390us	3.6390us	3.6390us	cudaDeviceGetPCIBusId
	0.00%	2.1940us	3	731ns	167ns	1.5810us	cudaDeviceGetCount
	0.00%	2.1750us	1	2.1750us	2.1750us	2.1750us	cudaEventSynchronize
	0.00%	1.7600us	1	1.7600us	1.7600us	1.7600us	cudaEventElapsedTime
	0.00%	1.0430us	2	521ns	290ns	753ns	cudaDeviceGet
	0.00%	287ns	1	287ns	287ns	287ns	cudaDeviceGetUuid

```

==384== Unified Memory profiling result:
Device "Tesla T4 (0)"
Count Avg Size Min Size Max Size Total Size Total Time Name
3 1.9076MB 1.7227MB 2.0000MB 5.722656MB 501.9200us Host To Device
12 163.00KB 4.0000KB 928.00KB 1.910156MB 173.5040us Device To Host
Total CPU Page faults: 24

```

Figure 5.14: Unified without prefetching 1000000

From Figure 5.9 to 5.11 and Figure 5.12 to 5.14, we can see that for the unified memory without the prefetching there will be host-to-device page fault generated for GPU but with the unified memory with prefetching there won't be any page fault and they both just call kernel once unlike the explicit CUDA code which will call kernel function several times. This actually makes sense since if we just use unified memory, managed memory may not be physically allocated when `cudaMallocManaged()` returns; it may only be populated on access (or prefetching) the memory entries may not be created until they are accessed by GPU or CPU and this may create page fault. After prefetching we can move the data to the GPU after initializing it and this will prevent the page fault.

Reference:

- [1] M. Harris. "Unified Memory for CUDA Beginners". NVIDIA Developer Blog. <https://developer.nvidia.com/blog/unified-memory-cuda-beginners/> (accessed Oct. 18, 2020).
- [2] D. Negrut, R. Serban, A. Li, A. Seidl. "Unified Memory in CUDA 6: A Brief Overview". Dr.Dobb's. <https://www.drdobbs.com/architecture-and-design/unified-memory-in-cuda-6-a-brief-overview/240169095?pgno=2> (accessed Oct. 18, 2020)